

```
# Install required packages via pip
!pip install spacy nltk torch torchcrf scikit-learn
```

 [Show hidden output](#)

```
# Download the spaCy English model
!pip python -m spacy download en_core_web_sm
```

 [Show hidden output](#)

```
!python -c "import nltk; nltk.download('treebank')"
```

 [Show hidden output](#)

```
!pip install pytorch-crf
```

 [Show hidden output](#)

```
import spacy
import numpy as np
import torch
import torch.nn as nn
import torch.optim as optim
from torchcrf import CRF
from nltk.corpus import treebank
from sklearn.model_selection import train_test_split
from torch.utils.data import DataLoader, Dataset
```

```
# Load spaCy model
nlp = spacy.load("en_core_web_sm")
```

```
# Load Treebank dataset
sentences = treebank.sents()
tags = [[tag for _, tag in sent] for sent in treebank.tagged_sents()]
```

```
# Build vocabulary
word2idx = {word: idx + 1 for idx, word in enumerate(set(word for sent in sentences for word in sent))}
tag2idx = {tag: idx for idx, tag in enumerate(set(tag for sent in tags for tag in sent))}
word2idx["<PAD>"] = 0
idx2tag = {idx: tag for tag, idx in tag2idx.items()}
```

```
# Convert data to indices
X = [[word2idx[word] for word in sent] for sent in sentences]
y = [[tag2idx[tag] for tag in sent] for sent in tags]
```

```
4
# Padding sequences
max_len = max(len(sent) for sent in X)
X = [sent + [0] * (max_len - len(sent)) for sent in X]
y = [sent + [tag2idx['NN']] * (max_len - len(sent)) for sent in y] # NN as default padding label
```

```
# Split dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Convert to PyTorch tensors
X_train, X_test = torch.tensor(X_train, dtype=torch.long), torch.tensor(X_test, dtype=torch.long)
y_train, y_test = torch.tensor(y_train, dtype=torch.long), torch.tensor(y_test, dtype=torch.long)
```

```
# Dataset class
class POSTaggingDataset(Dataset):
    def __init__(self, X, y):
        self.X = X
        self.y = y

    def __len__(self):
        return len(self.X)

    def __getitem__(self, idx):
        return self.X[idx], self.y[idx]

train_dataset = POSTaggingDataset(X_train, y_train)
test_dataset = POSTaggingDataset(X_test, y_test)
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)

# Bi-LSTM with CRF model
class BiLSTM_CRF(nn.Module):
    def __init__(self, vocab_size, tagset_size, embedding_dim=128, hidden_dim=256):
        super(BiLSTM_CRF, self).__init__()
        self.embedding = nn.Embedding(vocab_size, embedding_dim, padding_idx=0)
        self.lstm = nn.LSTM(embedding_dim, hidden_dim // 2, num_layers=2, bidirectional=True, batch_first=True)
        self.fc = nn.Linear(hidden_dim, tagset_size)
        self.crf = CRF(tagset_size, batch_first=True)

    def forward(self, x, tags=None):
        x = self.embedding(x)
        x, _ = self.lstm(x)
        emissions = self.fc(x)
        # Create a mask with the correct shape
        mask = (x[:, :, 0] != 0) # Assume padding token is 0 and is the first dimension in embedding

        if tags is not None:
            return -self.crf(emissions, tags, mask=mask, reduction='mean')
        else:
            return self.crf.decode(emissions, mask=mask) # Also apply mask during decoding

# Model, optimizer, loss
model = BiLSTM_CRF(len(word2idx), len(tag2idx))
optimizer = optim.Adam(model.parameters(), lr=0.001)


# Training
num_epochs = 10
for epoch in range(num_epochs):
    model.train()
    total_loss = 0
    for X_batch, y_batch in train_loader:
        optimizer.zero_grad()
        loss = model(X_batch, y_batch)
        loss.backward()
        optimizer.step()
        total_loss += loss.item()
    print(f"Epoch {epoch+1}/{num_epochs}, Loss: {total_loss / len(train_loader)}")
```

 Show hidden output

```
# Separate cell for evaluating accuracy on the test set
model.eval() # Set model to evaluation mode
total_correct = 0
total_tokens = 0

with torch.no_grad():
    for X_batch, y_batch in test_loader:
        predictions = model(X_batch) # Returns a list of lists with predicted tag indices
        for pred_seq, true_seq in zip(predictions, y_batch.tolist()):
            # Compare each token's predicted tag with the true tag
            for pred, true in zip(pred_seq, true_seq):
                total_correct += (pred == true)
                total_tokens += 1

accuracy = (total_correct / total_tokens)*100
print(f"Accuracy on test set: {accuracy:.4f}")
```

 Accuracy on test set: 99.0956

```

from sklearn.metrics import classification_report

# Collect true and predicted labels
true_labels = []
predicted_labels = []

model.eval()
with torch.no_grad():
    for X_batch, y_batch in test_loader:
        predictions = model(X_batch) # Get predicted indices
        for pred_seq, true_seq in zip(predictions, y_batch.tolist()):
            true_labels.extend(true_seq)
            predicted_labels.extend(pred_seq)

# Convert indices back to POS tags
true_labels = [idx2tag[idx] for idx in true_labels]
predicted_labels = [idx2tag[idx] for idx in predicted_labels]

# Compute and print F1-score, Precision, Recall
print(classification_report(true_labels, predicted_labels))

```

 [Show hidden output](#)

```
!pip install gradio
```

 Collecting gradio  
 Downloading gradio-5.22.0-py3-none-any.whl.metadata (16 kB)  
 Collecting aiofiles<24.0,>=22.0 (from gradio)  
 Downloading aiofiles-23.2.1-py3-none-any.whl.metadata (9.7 kB)  
 Requirement already satisfied: anyio<5.0,>=3.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (4.9.0)  
 Collecting fastapi<1.0,>=0.115.2 (from gradio)  
 Downloading fastapi-0.115.11-py3-none-any.whl.metadata (27 kB)  
 Collecting ffmpy (from gradio)  
 Downloading ffmpy-0.5.0-py3-none-any.whl.metadata (3.0 kB)  
 Collecting gradio-client==1.8.0 (from gradio)  
 Downloading gradio\_client-1.8.0-py3-none-any.whl.metadata (7.1 kB)  
 Collecting groovy~=0.1 (from gradio)  
 Downloading groovy-0.1.2-py3-none-any.whl.metadata (6.1 kB)  
 Requirement already satisfied: httpx>=0.24.1 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.28.1)  
 Requirement already satisfied: huggingface-hub>=0.28.1 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.29.3)  
 Requirement already satisfied: jinja2<4.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.1.6)  
 Requirement already satisfied: markupsafe<4.0,>=2.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.0.2)  
 Requirement already satisfied: numpy<3.0,>=1.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.0.2)  
 Requirement already satisfied: orjson~=3.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.10.15)  
 Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from gradio) (24.2)  
 Requirement already satisfied: pandas<3.0,>=1.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.2.2)  
 Requirement already satisfied: pillow<12.0,>=8.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (11.1.0)  
 Requirement already satisfied: pydantic>=2.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.10.6)  
 Collecting pydub (from gradio)  
 Downloading pydub-0.25.1-py2.py3-none-any.whl.metadata (1.4 kB)  
 Collecting python-multipart>=0.0.18 (from gradio)  
 Downloading python\_multipart-0.0.20-py3-none-any.whl.metadata (1.8 kB)  
 Requirement already satisfied: pyyaml<7.0,>=5.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (6.0.2)  
 Collecting ruff>=0.9.3 (from gradio)  
 Downloading ruff-0.11.1-py3-none-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl.metadata (25 kB)  
 Collecting safehttpx<0.2.0,>=0.1.6 (from gradio)  
 Downloading safehttpx-0.1.6-py3-none-any.whl.metadata (4.2 kB)  
 Collecting semantic-version~=2.0 (from gradio)  
 Downloading semantic\_version-2.10.0-py2.py3-none-any.whl.metadata (9.7 kB)  
 Collecting starlette<1.0,>=0.40.0 (from gradio)  
 Downloading starlette-0.46.1-py3-none-any.whl.metadata (6.2 kB)  
 Collecting tomlkit<0.14.0,>=0.12.0 (from gradio)  
 Downloading tomlkit-0.13.2-py3-none-any.whl.metadata (2.7 kB)  
 Requirement already satisfied: typer<1.0,>=0.12 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.15.2)  
 Requirement already satisfied: typing-extensions~=4.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (4.12.2)  
 Collecting uvicorn>=0.14.0 (from gradio)  
 Downloading uvicorn-0.34.0-py3-none-any.whl.metadata (6.5 kB)  
 Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages (from gradio-client==1.8.0->gradio) (2025.3.0)  
 Requirement already satisfied: websockets<16.0,>=10.0 in /usr/local/lib/python3.11/dist-packages (from gradio-client==1.8.0->gradio) (13.1)  
 Requirement already satisfied: idna>=2.8 in /usr/local/lib/python3.11/dist-packages (from anyio<5.0,>=3.0->gradio) (3.10)  
 Requirement already satisfied: sniffio>=1.1 in /usr/local/lib/python3.11/dist-packages (from anyio<5.0,>=3.0->gradio) (1.3.1)  
 Requirement already satisfied: certifi in /usr/local/lib/python3.11/dist-packages (from httpx>=0.24.1->gradio) (2025.1.31)  
 Requirement already satisfied: httpcore==1.\* in /usr/local/lib/python3.11/dist-packages (from httpx>=0.24.1->gradio) (1.0.7)  
 Requirement already satisfied: h11<0.15,>=0.13 in /usr/local/lib/python3.11/dist-packages (from httpcore==1.\*->httpx>=0.24.1->gradio) (0.14.0)  
 Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.28.1->gradio) (3.18.0)  
 Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.28.1->gradio) (2.32.3)  
 Requirement already satisfied: tqdm>=4.42.1 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.28.1->gradio) (4.67.1)  
 Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas<3.0,>=1.0->gradio) (2.9.0.post0)  
 Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas<3.0,>=1.0->gradio) (2025.1)  
 Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas<3.0,>=1.0->gradio) (2025.1)  
 Requirement already satisfied: annotated-types>=0.6.0 in /usr/local/lib/python3.11/dist-packages (from pydantic>=2.0->gradio) (0.7.0)  
 Requirement already satisfied: pydantic-core==2.27.2 in /usr/local/lib/python3.11/dist-packages (from pydantic>=2.0->gradio) (2.27.0)  
 Requirement already satisfied: click>=8.0.0 in /usr/local/lib/python3.11/dist-packages (from typer<1.0,>=0.12->gradio) (8.1.8)

```
import gradio as gr
import torch
import spacy

# Load spaCy model
nlp = spacy.load("en_core_web_sm")

# Assume word2idx, idx2tag, and model are already defined and loaded

def predict_pos_tags(user_input):
    # Tokenize the input
    doc = nlp(user_input)
    tokens = [token.text for token in doc]

    # Convert tokens to indices
    token_ids = [word2idx.get(token, 0) for token in tokens]

    # Create tensor with batch dimension
    input_tensor = torch.tensor([token_ids], dtype=torch.long)

    # Set model to evaluation mode and get predictions
    model.eval()
    with torch.no_grad():
        predictions = model(input_tensor)
    predicted_tag_indices = predictions[0]

    # Convert predicted indices back to POS tag labels
    predicted_tags = [idx2tag[idx] for idx in predicted_tag_indices]

    # Format the output
    result = "Token\tPredicted POS Tag\n"
    result += "\n".join([f"{token}\t{tag}" for token, tag in zip(tokens, predicted_tags)])
    return result

# Create Gradio UI
interface = gr.Interface(
    fn=predict_pos_tags,
    inputs=gr.Textbox(lines=2, placeholder="Enter your sentence here..."),
    outputs="text",
    title="POS Tagger",
    description="Enter a sentence to get POS tags predicted by the BiLSTM-CRF model"
)

# Launch Gradio app
interface.launch()
```

 Running Gradio in a Colab notebook requires sharing enabled. Automatically setting `share=True` (you can turn this off by setting `s`

Colab notebook detected. To show errors in colab notebook, set debug=True in launch()

\* Running on public URL: <https://34d0d689f153af688c.gradio.live>

This share link expires in 72 hours. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal in the worki



**No interface is running right now**